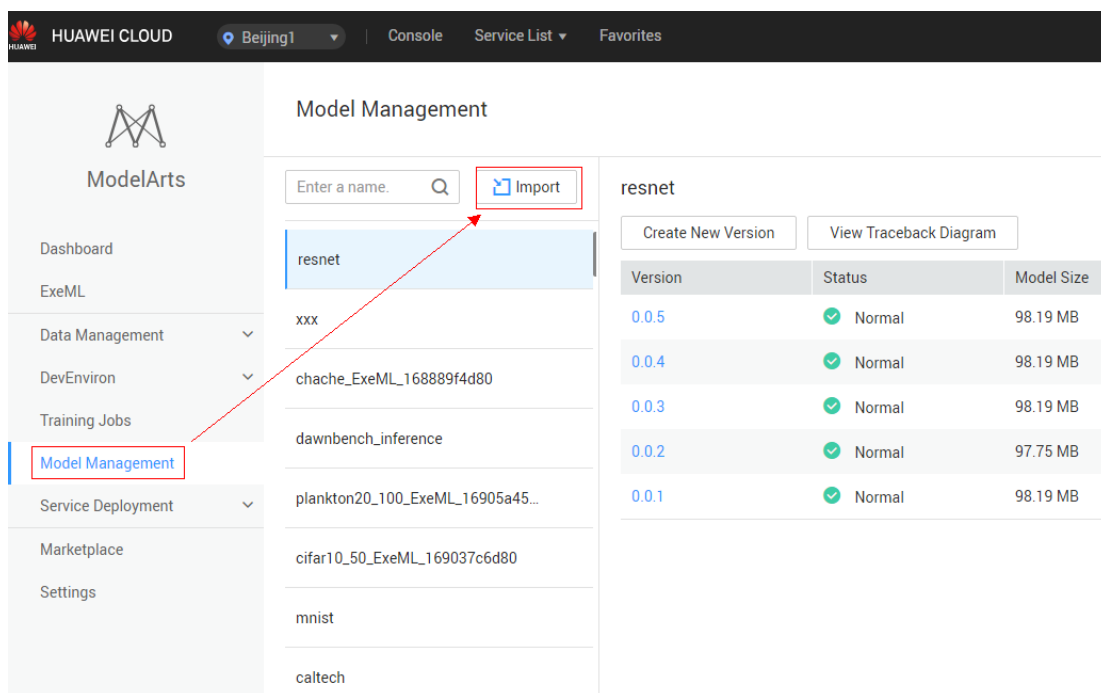


1. Deploy your model



Model Management

Enter a name.

resnet

xxx

chache_ExeML_168889f4d80

dawnbench_inference

plankton20_100_ExeML_16905a45...

cifar10_50_ExeML_169037c6d80

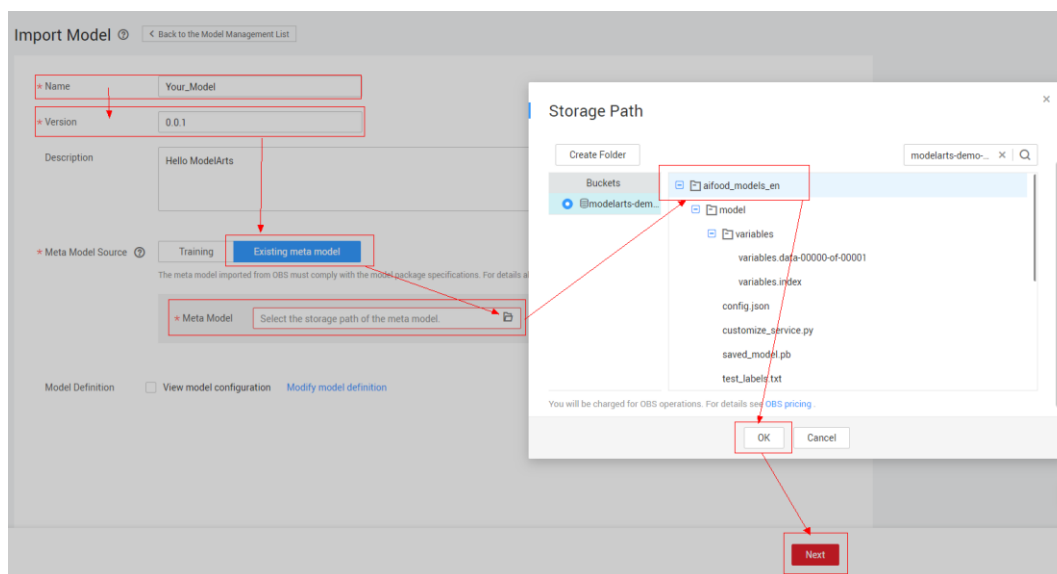
mnist

caltech

resnet

Create New Version View Traceback Diagram

Version	Status	Model Size
0.0.5	Normal	98.19 MB
0.0.4	Normal	98.19 MB
0.0.3	Normal	98.19 MB
0.0.2	Normal	97.75 MB
0.0.1	Normal	98.19 MB



Import Model

Name: Your_Model

Version: 0.0.1

Description: Hello ModelArts

Meta Model Source: Training

Meta Model: Existing meta model

Select the storage path of the meta model.

Storage Path

Create Folder

Buckets: @modelarts-dem...

aifood_models_en

model

variables

variables.data-00000-of-00001

variables.index

config.json

customize_service.py

saved_model.pb

test_labels.txt

OK Cancel

Next

1.1. Test your model through Batch Services

Deploy
Back to Batch Services

Basic Information

Name
My_Batch_Service

Description
Hello ModelArts

Parameter Configuration

Model and Version
Your_Model
0.0.1

Input Path
/obs-dataset/food101/food-101-small-sample/judge_samples/

Request Path
/

Mapping Type
CSV
File

Environment Variable
Add Environment Variable

Output Path
/obs-dataset/food101/food-101-pred-005/

Instance Flavor
2core | 8GiB
2core | 8GiB | 1*P4

Instance Count
1

Price Free for a limited time
This price is an estimate and may differ from the final price. Pricing details

Next

If something goes wrong with your model, you can find error message in “Log”

ModelArts

Batch Service - My_Batch_Service

Dashboard

EveML

Data Management

DevEnviron

Training Jobs

Model Management

Service Deployment

Marketplace

Settings

Batch Services

Log

Name
My_Batch_Service

Status
Running

Calculate node flavor
2core | 8GiB | 1*P4

Calculate node num
1

Description
Hello ModelArts

Service ID
2725a0f1-22c7-4953-9b63-622a18833ff4

Model Name & Version
Your_Model 0.0.1

Input Path
https://obs-dataset-obs.myhwclouds.com/food101/food-101-small-sample/judge_samples/

Environment Variable
-

Output Path
https://obs-dataset-obs.myhwclouds.com/food101/food-101-pred-005/

Log

Traceback Diagrams

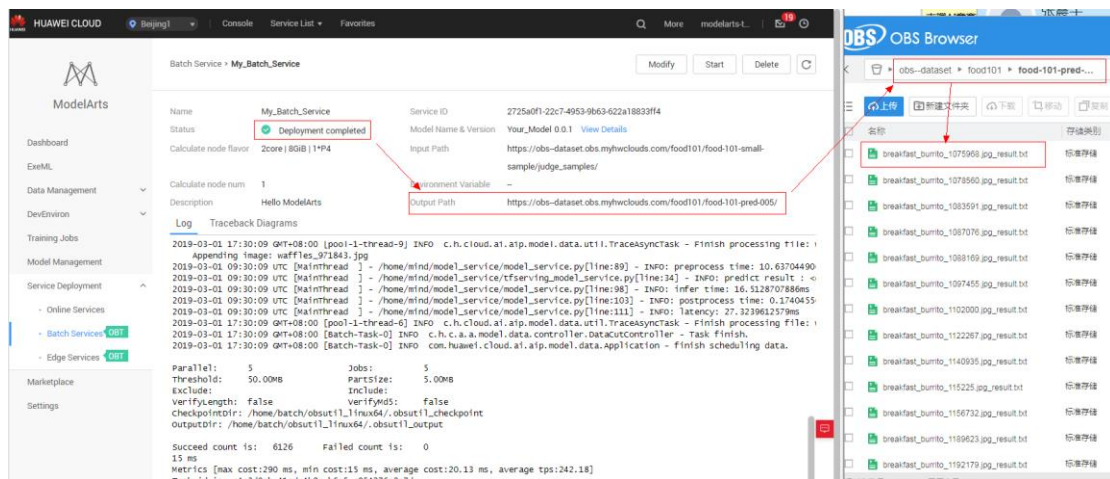
```

2019-03-01 17:28:17 UTC [pool-1-thread-4] INFO c.h.cloud.ai.alp.model.data.util.TraceAsyncTask - Begin processing file: canno11_3085374.jpg
Appending Image: canno11_3085374.jpg
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:89] - INFO: preprocess time: 61.4500043776ms
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:89] - INFO: preprocess time: 10.2391242981ms
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/tfserving_model_service.py[Line:34] - INFO: predict result: <google.protobuf.pyext._message.MessageAppContainer object at 0x...
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:98] - INFO: infer time: 17.6420211792ms
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:103] - INFO: postprocess time: 0.192880630493ms
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:111] - INFO: latency: 79.2849063873ms
2019-03-01 17:28:17 UTC [Batch-Task-0] INFO c.h.cloud.ai.alp.model.data.util.TraceAsyncTask - Finish processing file: canno11_3086498.jpg
2019-03-01 17:28:17 UTC [Batch-Task-0] INFO c.h.cloud.ai.alp.model.data.util.TraceAsyncTask - Begin processing file: canno11_3285277.jpg
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/tfserving_model_service.py[Line:34] - INFO: predict result: <google.protobuf.pyext._message.MessageAppContainer object at 0x...
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:98] - INFO: infer time: 16.891002655ms
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:103] - INFO: postprocess time: 0.165939331059ms
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:111] - INFO: latency: 27.2960666842ms
2019-03-01 17:28:17 UTC [pool-1-thread-3] INFO c.h.cloud.ai.alp.model.data.util.TraceAsyncTask - Finish processing file: canno11_3825694.jpg
2019-03-01 17:28:17 UTC [pool-1-thread-3] INFO c.h.cloud.ai.alp.model.data.util.TraceAsyncTask - Begin processing file: canno11_3106000.jpg
Appending Image: canno11_2866470.jpg
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:89] - INFO: preprocess time: 13.6239528656ms
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:89] - INFO: preprocess time: 11.3429691315ms
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/tfserving_model_service.py[Line:34] - INFO: predict result: <google.protobuf.pyext._message.MessageAppContainer object at 0x...
2019-03-01 09:28:17 UTC [mainthread ] - /home/mind/model_service/model_service.py[Line:98] - INFO: infer time: 17.0619447762ms

```

Download a result file to see if it meets the specifications defined by scoring system.

2019-3-1

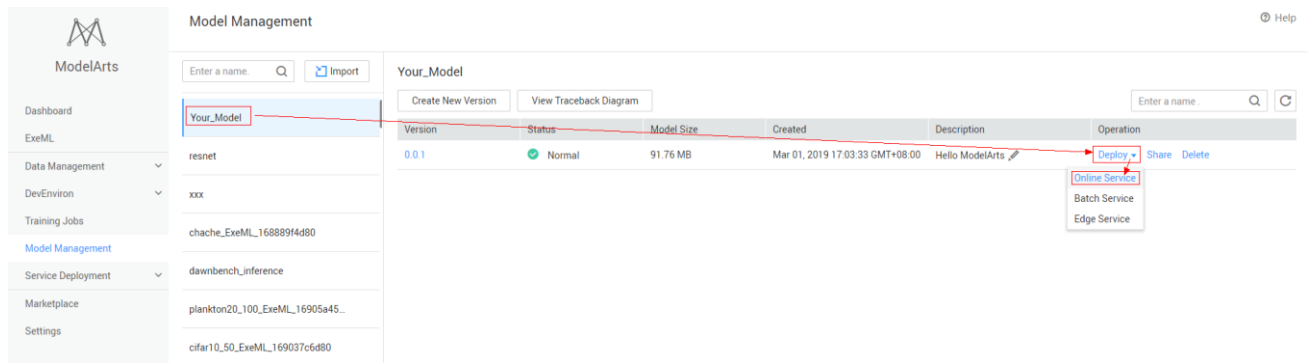


The screenshot displays the Huawei Cloud ModelArts console interface. On the left, the 'ModelArts' sidebar is visible with options like Dashboard, Data Management, DevEnviron, Training Jobs, Model Management, and Service Deployment. The main area shows the 'Batch Service' details for 'My_Batch_Service'. Key information includes:

- Status:** Deployment completed
- Service ID:** 2725a0f1-22c7-4953-9b63-422a18833f4
- Model Name & Version:** Your_Model 0.0.1
- Input Path:** https://obs-dataset.obs.myhuaweiclouds.com/food101/food-101-small-sample/judge_samples/
- Output Path:** https://obs-dataset.obs.myhuaweiclouds.com/food101/food-101-pred-005/

 The 'Log' section shows a series of INFO messages from the 'TraceSyncTask' and 'DataCutController' indicating successful processing and scheduling. On the right, the 'OBS Browser' window shows a directory structure with files like 'breakfast_burnto_1075968.jpg_result.txt'.

1.2. Test your model through Online Service



The screenshot shows the 'Model Management' section of the Huawei Cloud ModelArts console. The 'Your_Model' table lists the following models:

Version	Status	Model Size	Created	Description	Operation
0.0.1	Normal	91.76 MB	Mar 01, 2019 17:03:33 GMT+08:00	Hello ModelArts	Deploy, Share, Delete

 The 'Deploy' button is highlighted, and a dropdown menu shows options: 'Online Service', 'Batch Service', and 'Edge Service'. The 'Online Service' option is selected. The left sidebar shows the 'Model Management' menu item.

Deploy
Back to Online Services

Basic Information

* Name
My_Online_Service

Description
Hello ModelArts
15/100

* Billing Mode
Pay-per-use

Network Configuration

Custom Configuration
?

Parameter Configuration

* Model and Configuration

Model Management
Your_Model
0.0.1
Instance Flavor
2core | 8GiB

Traffic Ratio (%)
100
Instance Count
1

Price ¥ 0.80/hour
This price is an estimate and may differ from the final price. [Pricing details](#)

Next

After Online Service is ready


ModelArts
Dashboard
ExamLE
Data Management
DevEnviron
Training Jobs
Model Management
Service Deployment
Online Services
Batch Services
Edge Services
Marketplace
Settings

My_Online_Service
Running
Service ID: 81854f63-9384-411a-b8fb-b6452c160afe
Type: Self-created service

Description: Hello ModelArts

Usage Guides
Configurations
Configuration Update Record
Tests
Monitoring Data
Log Updates
Sharing
Traceback Diagrams

Code Debugging
File



Response

```

1 {
2   "logits": {
3     "cheese": 1.524839520454067,
4     "str_fried_tortoise": -1.7106966972351074,
5     "glutinous_rice_ball_with_peanut_and_sesame_sweets": 0.2637473,
6     "samban_porridge": 0.88605582241638,
7     "glutinous_rice_dumpling": 4.4460238925390625,
8     "layer_cake": -0.7078758478164673,
9     "green_bean_cake": -1.371852159500122,
10    "hot_dog": -0.4825528264045713,
11    "pop_corn": -2.7938296794891357,
12    "persimmon_pancake": -2.035639762878418,
13    "omelette": 0.035944193601680276,
14    "poon_choi": -1.2870237827301025,
15    "claypot": -0.41622111201286316,
16    "black_sesame_paste": -0.4339074263188507,
17    "frozen_lemon_tea": -0.9952121891975403,
18    "onion_rings": -2.330695152282715,
19    "fried_chestnuts": -1.6958777304510498,
20    "mango_pome_lo_sago": 0.34968435764312744,
21    "french_fries": 1.55490825604246

```

Test File: Chrysanthemum... (508.70KB)
Predicted Success

If something goes wrong with your model, you can find error message in "Log Updates"



ModelArts

Dashboard

ExeML

Data Management

DevEnvision

Training Jobs

Model Management

Service Deployment

- Online Services
- Batch Services
- Edge Services

Marketplace

Settings

Online Service • My_Online_Service

NameMy_Online_ServiceService ID81854f53-938d-411a-bfdb-b6452c160afe

StatusRunningTypeSelf-created service

Failed Calls/Total Calls0/1Network Configuration-

DescriptionHello ModelArts

Usage Guides

Configurations

Configuration Update Record

Tests

Monitoring Data

Log Updates

Sharing

Traceback Diagrams

Your_Model_0.0.1

All Nodes

Latest 5 minutes

Latest 30 minutes

Latest 1 hour

Custom

Support exact search and fuzzy

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_def/model_def.py[111:84] - INFO: model: 3 signed int, 16 float, 1 object

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_def/model_def.py[111:42] - INFO: model inputs: {'images': tf.float32}

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_def/model_def.py[111:53] - INFO: model outputs: {'logits': tf.float32}

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_def/model_def.py[111:55] - INFO: model path: /home/mind/model/1

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_def/model_def.py[111:24] - INFO: signature name: predict_object

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_def/model_def.py[111:31] - INFO: model signature: predict_object

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_def/model_def.py[111:42] - INFO: model inputs: {'images': tf.float32}

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_def/model_def.py[111:53] - INFO: model outputs: {'logits': tf.float32}

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_def/model_def.py[111:55] - INFO: model path: /home/mind/model/1

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:234] - INFO: connecting to TensorFlow server 127.0.0.1:8500

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:234] - INFO: connecting to TensorFlow server 127.0.0.1:8500

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:89] - INFO: preprocess time: 102.49900817ms

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:98] - INFO: infer time: 7039.6828654ms

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:98] - INFO: infer time: 7039.6828654ms

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:103] - INFO: postprocess time: 18.6021327972ms

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:111] - INFO: latency: 7160.78400612ms

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:89] - INFO: preprocess time: 102.49900817ms

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:98] - INFO: infer time: 7039.6828654ms

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:98] - INFO: infer time: 7039.6828654ms

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:103] - INFO: postprocess time: 18.6021327972ms

2019-03-01 09:14:04 UTC [kainthread] - /home/mind/model_service/model_service.py[111:111] - INFO: latency: 7160.78400612ms

2019-3-1